



a. F1 vs. Training Data Fraction



b. CIs at 25% Training



Dashed line in plot (a) shows GPT-4.1 mini few-shot baseline.

95% confidence intervals for macro F1 scores in plot (b) calculated via bootstrap resampling.