

2026-04-08 Prepared for the APSA Virtual Research Meeting.
Access the code, data, and analysis at <https://github.com/eteitelbaum/code-satp.git>

Extracting Event Counts and Locations from Conflict Reports^{*}

Evidence from the South Asia Terrorism Portal

Emmanuel Teitelbaum[†] 
George Washington University
Department of Political Science
ejt@gwu.edu

Funding: This project received funding from the Smith Richardson Foundation and in-kind support from the GW Institute for International Economic Policy.

^{*}I am grateful to Mohiddin Bacha Shaik and Shiva Sharma for their excellent research assistance.

[†]Corresponding author.

Table 1: Death count extraction performance: seq2seq models.

Metric	Conflibert	T5-Base	T5-Large	IndicBART	mT5-Base	NT5-Small	T5-XL
MAE	0.37	0.27	0.16†	0.33	0.27	0.29	0.12
RMSE	2.21	3.21	1.18†	1.73	1.61	3.11	0.85
Within-1	94.6	97.7†	97.8†	94.9	96.3	97.3†	98.1
Within-2	97.3	98.6†	98.7†	96.5	97.7	98.4†	98.7
Nonzero MAE	0.56	0.39	0.21†	0.42	0.30	0.39	0.15

Bold = significantly better than Conflibert (paired bootstrap, $n = 5000$, $p < .05$). † = not significantly different from ceiling ($p > .05$). MAE and Nonzero MAE: lower is better. Within-1 and Within-2: higher is better.

Table 2: Death count extraction performance: large language models.

Metric	Conflibert	Llama-3.1-8B	Mistral-7B	Mixtral-8x7B	GPT-4o-mini
MAE	0.37	0.23†	0.30	0.30†	0.13
RMSE	2.21	2.77†	3.04	3.75†	0.94
Within-1	94.6	97.7†	97.2†	97.9†	97.9
Within-2	97.3	98.4†	97.8	98.5†	98.6
Nonzero MAE	0.56	0.19†	0.22†	0.16†	0.17

Bold = significantly better than Conflibert (paired bootstrap, $n = 5000$, $p < .05$). † = not significantly different from ceiling ($p > .05$). MAE and Nonzero MAE: lower is better. Within-1 and Within-2: higher is better.

Introduction

[Introduction text to be written.]

Data and Methods

[Methods text to be written.]

Results

Baseline Model Comparisons

[Results text to be written. Tables Table 1 and Table 2 report overall performance for the seq2seq and LLM model families respectively. Values in bold are significantly better than the Conflibert-Poisson baseline (paired bootstrap, $n = 5,000$ replications, $p < .05$); a dagger (†) indicates the model is not significantly different from the ceiling model.]

Bin-Level Performance

?@fig-heatmap-seq2seq

Rare-Bin Improvement Strategies

Seq2Seq Training-Side Interventions

Table 3

LLM Prompt Engineering Interventions

Table 4

Discussion

[Discussion text to be written.]




Figure 1: Exact match (%) by death count bin. Seq2seq models (top) and LLMs (bottom), each ordered by overall MAE (best at top). Conflibert-Poisson appears in both panels as a reference baseline.

Table 3: Seq2seq rare-bin intervention results (Flan-T5-Large base).

Metric	Baseline	Weighted	Loss Wt.	Targeted	Back-transl.	Paraphrase	T5-XL
MAE	0.16†	0.26	0.22†	0.24†	0.27	0.27	0.12
Nonzero MAE	0.21†	0.36	0.35	0.35	0.37	0.38	0.15
Bin 3-5 MAE	0.24†	0.27	0.24	0.24	0.21†	0.23†	0.19
Bin 3-5 Exact (%)	86.7†	85.3	86.7†	85.3	88.0†	86.7†	90.7
Bin 6+ MAE	1.97†	4.08	3.95	4.05	4.18	4.34	1.05
Bin 6+ Exact (%)	71.1†	73.7†	68.4	68.4	71.1	71.1	84.2

Bold = significantly better than baseline model (paired bootstrap, $n = 5000$, $p < .05$). † = not significantly different from ceiling ($p > .05$). MAE rows: lower is better. Exact (%) rows: higher is better.

Table 4: LLM prompt engineering results (Llama-3.1-8B base).

Metric	Baseline	Protocol	Bin-shot	Hard-shot	Combined	GPT-4o-mini
MAE	0.23†	0.32	0.20†	0.11†	0.11†	0.13
Nonzero MAE	0.19†	0.20†	0.15†	0.14†	0.13†	0.17
Bin 3-5 MAE	0.20†	0.23†	0.28	0.20†	0.28	0.16
Bin 3-5 Exact (%)	90.7†	88.0†	88.0	90.7†	88.0	93.3
Bin 6+ MAE	1.53†	0.95†	0.97†	0.84†	0.63†	1.24
Bin 6+ Exact (%)	78.9†	76.3†	78.9†	81.6†	89.5†	84.2

Bold = significantly better than baseline model (paired bootstrap, $n = 5000$, $p < .05$). † = not significantly different from ceiling ($p > .05$). MAE rows: lower is better. Exact (%) rows: higher is better.

Appendix

References